



From Raw Sensor Data to Semantic Web Triples

Information Flow in Semantic Sensor Networks

Nikolaos Konstantinou, Ph.D.
Autonomic and Grid Computing group
Athens Information Technology



Lecture Outline

- Introduction
- Sensor data
- Semantic web
- Context-awareness
- The GSN middleware
- Exposing sensor data as triples
- Semantic Integration using SPARQL



Introduction

- The Situation
 - Decrease in the value of sensors encourage the shift from Desktop to Ubiquitous Computing
- The Problem
 - Extracting meaning from a network of deployed sensors
- Why is this a problem?
 - Raw sensor data is useless unless properly annotated
 - Limited resources in terms of processing, storage capabilities and bandwidth
- What is the suggested solution?
 - Establish ways to automatically process and manage the data



Concepts involved (1)

- Context-awareness
 - Context-aware systems are able to sense and measure their environment and include these measurements in their behavior
- Data fusion
 - Combine information and data residing at disparate sources, in order to achieve improved accuracies and more specific inferences than could be achieved by the use of a single data source alone
 - Data fusion spans various levels, from signal and object refinement (low level) to situation and threat assessment (high level)



Concepts involved (2)

- Information Integration
 - Unify information originating from various sources in order to allow its processing as a whole
 - Obstacles include
 - Heterogeneity in source schemas and data
 - Various Technical Spaces
 - Semantic and syntactic differences
- *Semantic* Information Integration
 - The resulting integration scheme carries its semantics
- Information Merging
 - Unification of the information at the implementation level



Concepts involved (3)

- Information Aggregation
 - Report the mean value of a set of measurements
 - E.g. the average temperature of a set of temperature sensors
- Information Annotation
 - Inclusion of metadata next or in the actual data
 - E.g. ID3 tags in mp3
- *Semantic* Information Annotation
 - Unambiguously define information
 - Third parties can understand the information



Lecture Outline

- Introduction
- Sensor data
- Semantic web
- Context-awareness
- The GSN middleware
- Exposing sensor data as triples
- Semantic integration using SPARQL



Sensors and Sensor Data (1)

- Sensors are devices that measure physical properties
 - Temperature, motion, light, humidity sensors
 - Also cameras, microphones, GPS-enabled smartphones
- Sensors provide data that can be
 - Streamed data
 - Audio/Video content
 - Event-based
 - Temperature measurement
 - RFID tag read
 - Light curtain interrupt




Sensors and Sensor Data (2)

- Why is sensor data any different than other forms of data, e.g. multimedia?
 - Synchronization issues
 - Apply acceptance thresholds
 - Erroneous measurements
 - Apply aggregation
 - Limited resources in the nodes
 - Caution when designing *where* the actual processing takes place
 - Keep a sliding window
 - Heavy process in the Gateway Nodes
 - Streaming may lead to packet losses
 - Reconstruct, or take decisions based on what you have

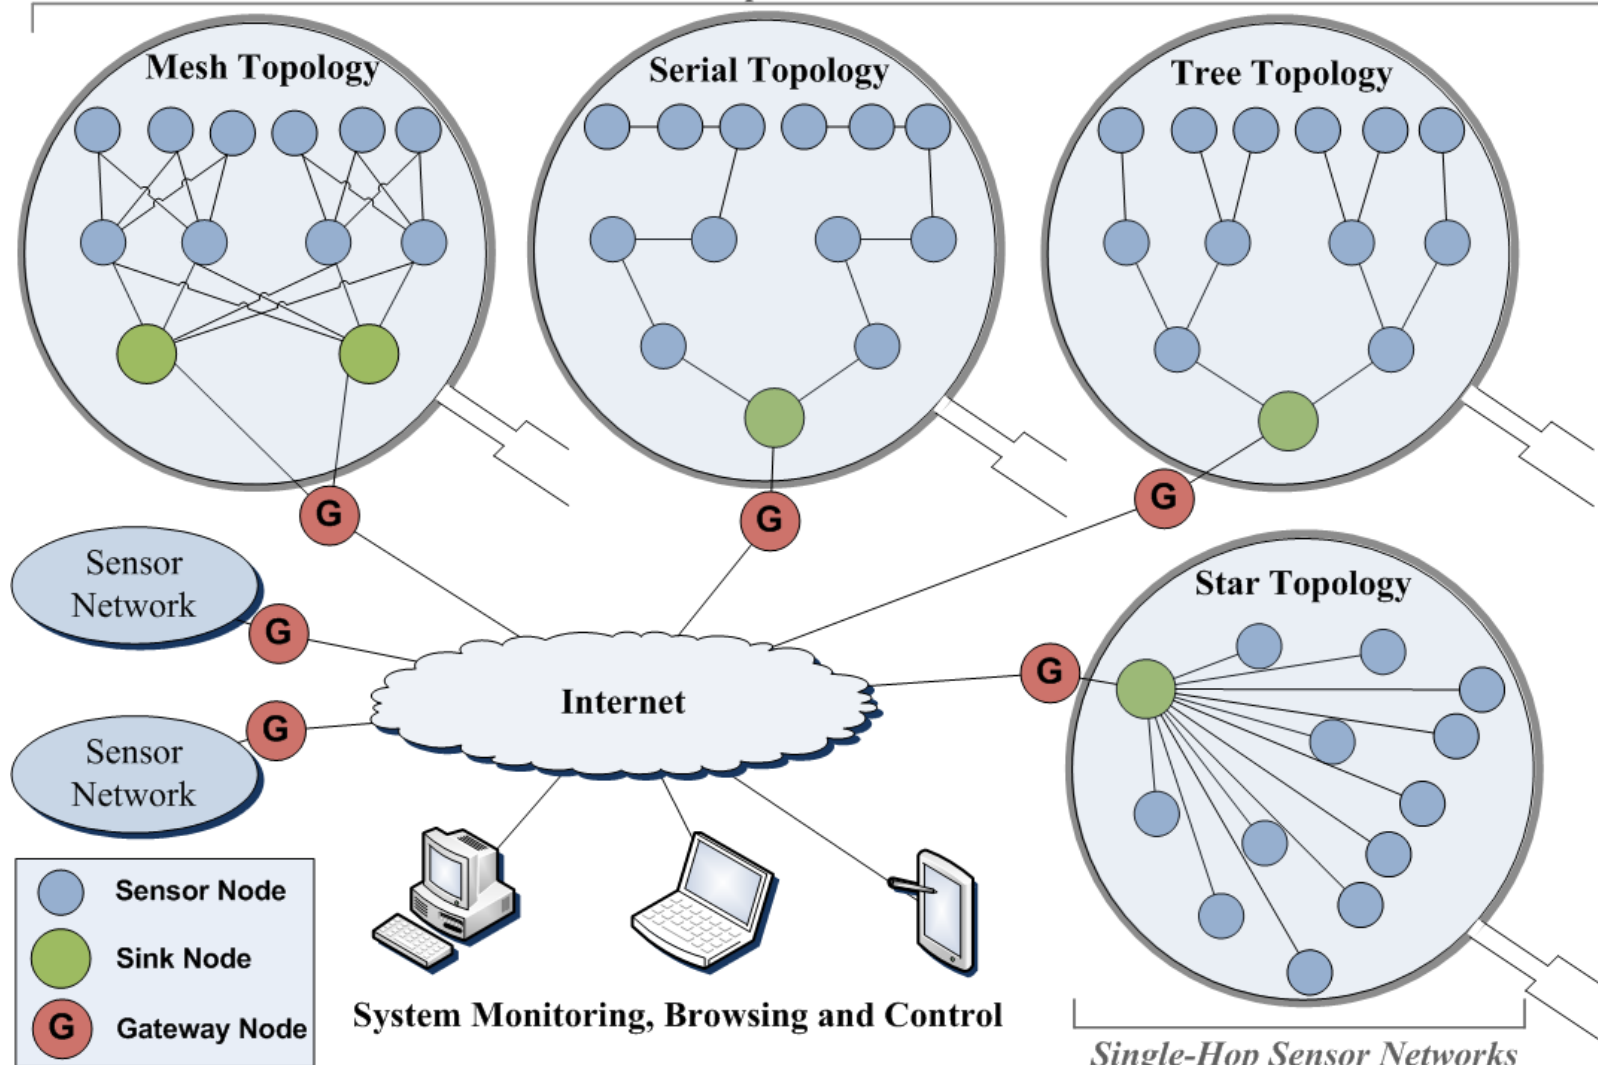
Sensor Network Topologies (1)

- Sink nodes collect information
 - Higher processing capabilities



Sensor Network Topologies (2)

Multi-Hop Sensor Networks





Lecture Outline

- Introduction
- Sensor Data
- Semantic Web
- Context-Awareness
- The GSN Middleware
- Exposing Sensor Data as triples
- Semantic Integration using SPARQL

Why Semantic Web (1)

- Knowledge in the form of a graph
 - (subject, property, object)
- Information is assigned an unambiguously defined meaning, its semantics
 - Queries can be posed by any third parties
 - Ontology, a well defined vocabulary
- Numerous ontologies already available and interconnected on the Web
 - Can and should be used when integration is a goal

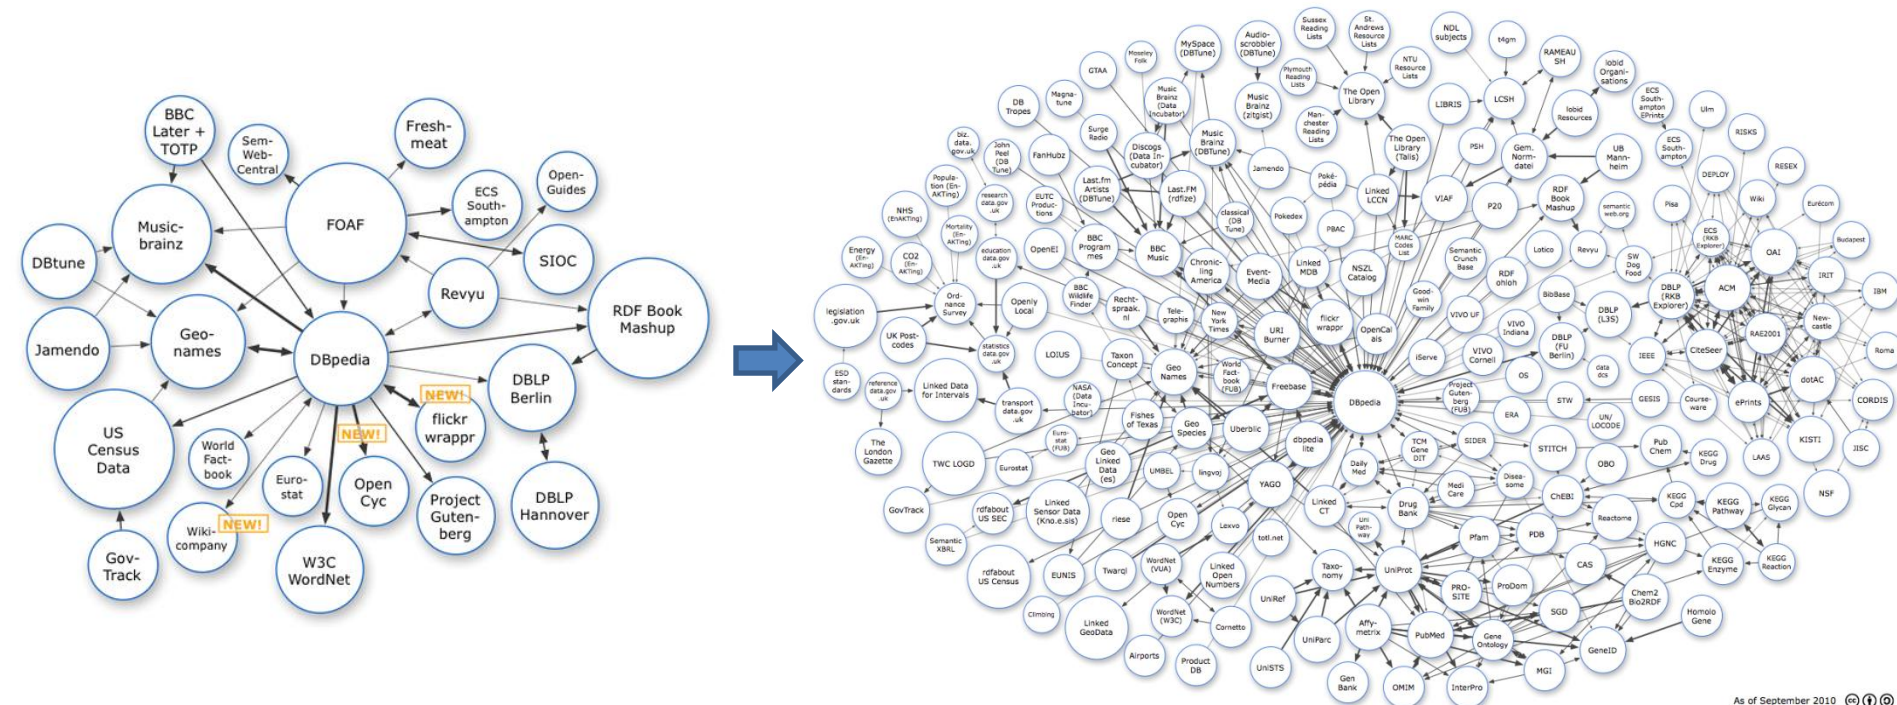




Why Semantic Web (2)

- Enables semantic annotation and integration
- Enables reasoning
 - Extract implicit information from the explicitly asserted
 - Assure concept consistency and satisfiability
- Open source tools available
 - Protégé, Jena, Virtuoso, ...
- Allows information to be exposed as Linked Open Data

The Linked Open Data Cloud



2007

2010


Source: linkeddata.org



Lecture Outline


- Introduction
- Sensor data
- Semantic web
- Context-awareness
- The GSN middleware
- Exposing sensor data as triples
- Semantic integration using SPARQL

Context-awareness: The big picture (1)



- The sensors capture information from the environment
- The signal processing components produce structured information
 - Face detector, Body Tracker, Vehicle tracker, Smoke detection etc.


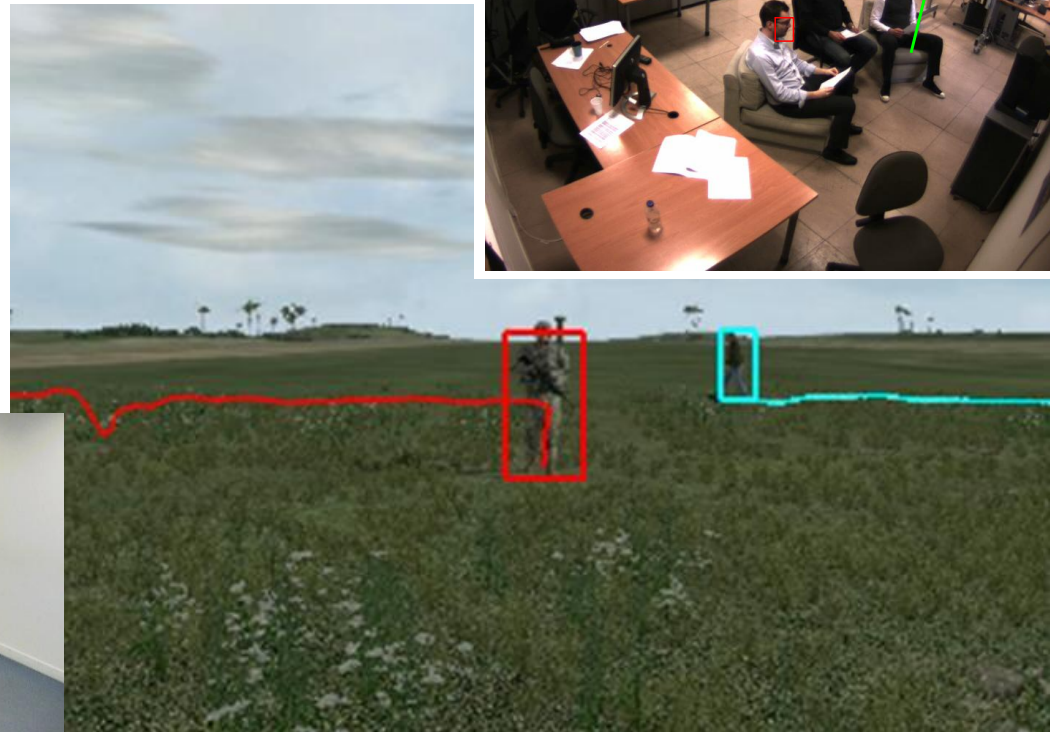

Context-awareness: The big picture (2)



- The middleware
 - Business logic: Program, configure, monitor and control system behavior
 - Storage layer
 - Support database: limited historicity, sliding window
 - Archive database: enables further processing
- Note the Information/Control duality

Signal Processing Components (1)

- Much work carried out in the AGC lab
 - Image Processing
 - Face Detection/Recognition/Tracking
 - Body Tracking
 - Audio Processing
 - A/V Localization
 - Voice activity detection





Signal Processing Components (2)

- Challenges
 - Processing is resource-hungry
 - Heterogeneous technologies must be combined
 - Production Algorithms in C++, prototypes in Matlab
 - Multidisciplinary skills required
 - Well-defined Interfaces need to be developed using RMI/Sockets/Web services/JNA
 - Video processing differs from streaming processing
 - Processing an avi file differs (greatly!) from processing an rtp stream



Lecture Outline

- Introduction
- Sensor data
- Semantic web
- Context-awareness
- The GSN middleware
- Exposing sensor data as triples
- Semantic integration using SPARQL



Global Sensor Networks (1)

- Open-source, java-based *middleware* solution
- Available online at sf.net/projects/gsn/
- Adaptability
 - Everything is a virtual sensor
 - Virtual sensors rely on wrappers
 - Every data producer can be integrated into the GSN with a virtual sensor and wrapper

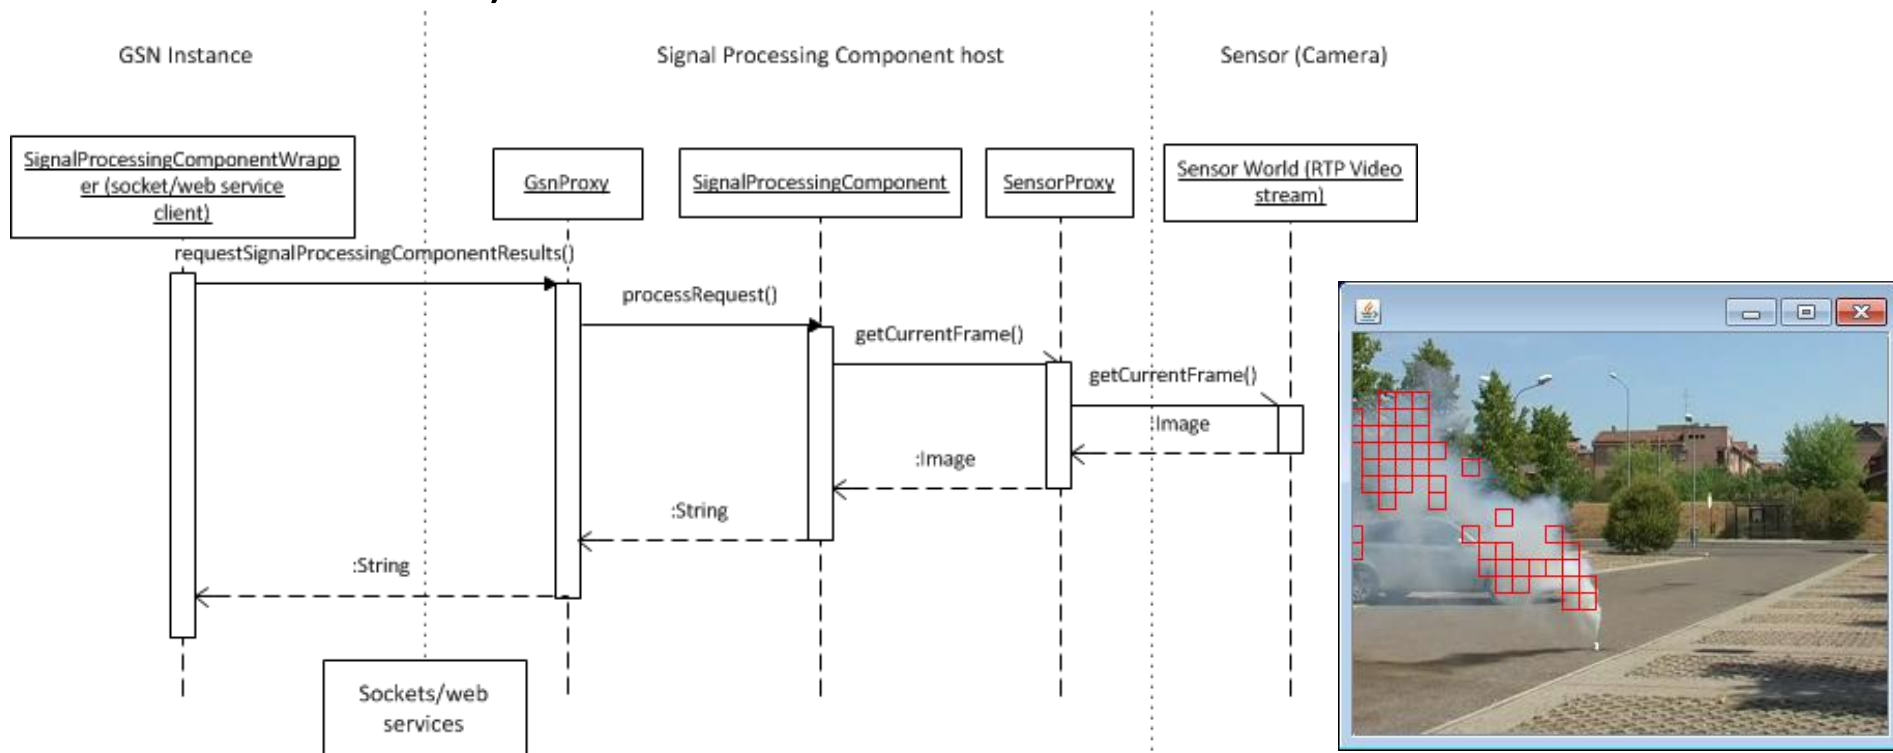
Global Sensor Networks (2)

- Simplicity
 - Configurable without compiling source code
 - Web application for sensor management
- Scalability
 - Allows communication between nodes
 - Allows data aggregation and fusion using an SQL-like declarative language



Global Sensor Networks (3)

- Example
 - Integrating a Signal Processing component (e.g. a Smoke detector) into GSN





Example of data aggregation using GSN

```
<output-structure>
```

```
...
```

```
<field name="TEMPERATURE" type="int" />
```

```
...
```

```
</output-structure>
```

```
...
```

```
<storage history-size="24h" />
```

```
...
```

```
<streams>
```

```
<stream name="input1">
```

```
<source alias="source1" sampling-rate="1" storage-size="1">
```

```
<address wrapper="temperature">
```

```
<predicate key="sampling-rate">10000</predicate>
```

```
</address>
```

```
<query>select TEMPERATURE from wrapper</query>
```

```
</source>
```

```
<query>select avg(TEMPERATURE) from source1</query>
```

```
</stream>
```

```
</streams>
```

measured
properties

sliding
window
size

data source

aggregated
output



Example of data fusion using GSN

```
<storage history-size="1h" />
```

```
<streams>
```

```
  <stream name="teststream" rate="1000">
```

```
    <source name="source1" alias="source1" storage-size="100" slide="0" sampling-rate="1">
```

```
      <address wrapper="remote-rest">
```

```
        <predicate key="HOST">localhost</predicate><predicate key="PORT">22001</predicate>
```

```
        <predicate key="QUERY">select FACE_COUNT from doorwatcher</predicate>
```

```
      </address>
```

```
      <query>select FACE_COUNT AS S1 from wrapper</query>
```

```
    </source>
```

```
    <source name="source2" alias="source2" storage-size="100" slide="0" sampling-rate="1">
```

```
      <address wrapper="remote-rest">
```

```
        <predicate key="HOST">localhost</predicate><predicate key="PORT">22002</predicate>
```

```
        <predicate key="QUERY">select TAG from touchatag</predicate>
```

```
      </address>
```

```
      <query>select TAG AS S2 from wrapper</query>
```

```
    </source>
```

```
  <query>
```

```
    select source1.S1 as S1OUT, source2.S2 as S2OUT
```

```
    from source1, source2
```

```
    where source1.S1 > 0 AND source2.S2="04dddc9232580"
```

```
  </query>
```

```
</stream>
```

```
</streams>
```

source 1

source 2

fused
output



The Storage Layer (1)

- Schema according to the Virtual Sensors
 - Virtual Sensor definition example:

```
<virtual-sensor name="temperature" priority="10">
```

```
...
```

```
<output-structure>
```

```
<field name="ID" type="int" />
```

```
<field name="SENSORTIME" type="time" />
```

```
<field name="TEMPERATURE" type="double" />
```

```
<field name="UNIT" type="varchar(255)" />
```

```
</output-structure>
```



The Storage Layer (2)

- Schema according to the Virtual Sensors
 - Schema auto-generated SQL create statement:

```
CREATE TABLE `gsn1`.`temperature` (
```

```
  `PK` bigint(20) NOT NULL AUTO_INCREMENT,
```

```
  `timed` bigint(20) NOT NULL,
```

```
  `ID` int(11) DEFAULT NULL,
```

```
  `SENSORTIME` time DEFAULT NULL,
```

```
  `TEMPERATURE` double DEFAULT NULL,
```

```
  `UNIT` varchar(255) DEFAULT NULL,
```

```
  PRIMARY KEY (`PK`),
```

```
  UNIQUE KEY `temperature_INDEX` (`timed`)
```

```
) ENGINE=MyISAM AUTO_INCREMENT=122 DEFAULT CHARSET=latin1
```





} added by GSN

} depending on the virtual sensor description

The Storage Layer (3)

- Historical data
 - According to the Virtual Sensor definition
 - Tuple-based: `<storage history-size="1" />`
 - Time-based: `<storage history-size="1m" />`
- Out-of-the-box support for
 - Mysql
 - SQL Server
 - Oracle
 - H2
 - Can be extended to support other RDBMS's

gsn1.temperature: 10 rows total (approximately)

 PK	 timed	 ID	TEMPERATURE	 UNIT
14983	1301312947745	3612	23	C
15023	1301312959138	3612	23	C
15018	1301312957746	3612	23	C
15013	1301312956379	3612	23	C
15008	1301312954970	3612	23	C
15003	1301312953579	3612	23	C
14998	1301312952170	3612	23	C
14993	1301312950700	3612	23	C
14988	1301312949223	3612	23	C
15028	1301312960725	3612	23	C



The Storage Layer (4)

- Storage Layer can be
 - Centralized (in the Central Control node)
 - Data is pushed to the Central Control node
 - Distributed (in the Gateway nodes)
- Interfaces
 - Legacy (ODBC/JDBC)
 - Web Services (SOAP or RESTful)
 - SPARQL Endpoints
 - Allow Semantic Information Integration



Lecture Outline

- Introduction
- Sensor data
- Semantic web
- Context-awareness
- The GSN middleware
- Exposing sensor data as triples
- Semantic integration using SPARQL



Exposing sensor data as triples (1)

- OpenLink Virtuoso universal server can be used as
 - A web application server
 - A relational database repository
 - A triplestore
 - A web service server
- Open-source version available at <http://virtuoso.openlinksw.com/>
- Cluster Configuration
 - Parallel and Horizontal scaling



Exposing sensor data as triples (2)

- Virtuoso RDF Views
 - Export relational data as triples
- SPARQL 1.1 support, plus
 - Full Text Queries
 - Geo Spatial Queries
 - Business Analytics and Intelligence
 - SQL Stored Procedure and Built-In Function exploitation from SPARQL
 - Create, Update, and Delete (SPARUL)
- Backward-chaining OWL reasoner

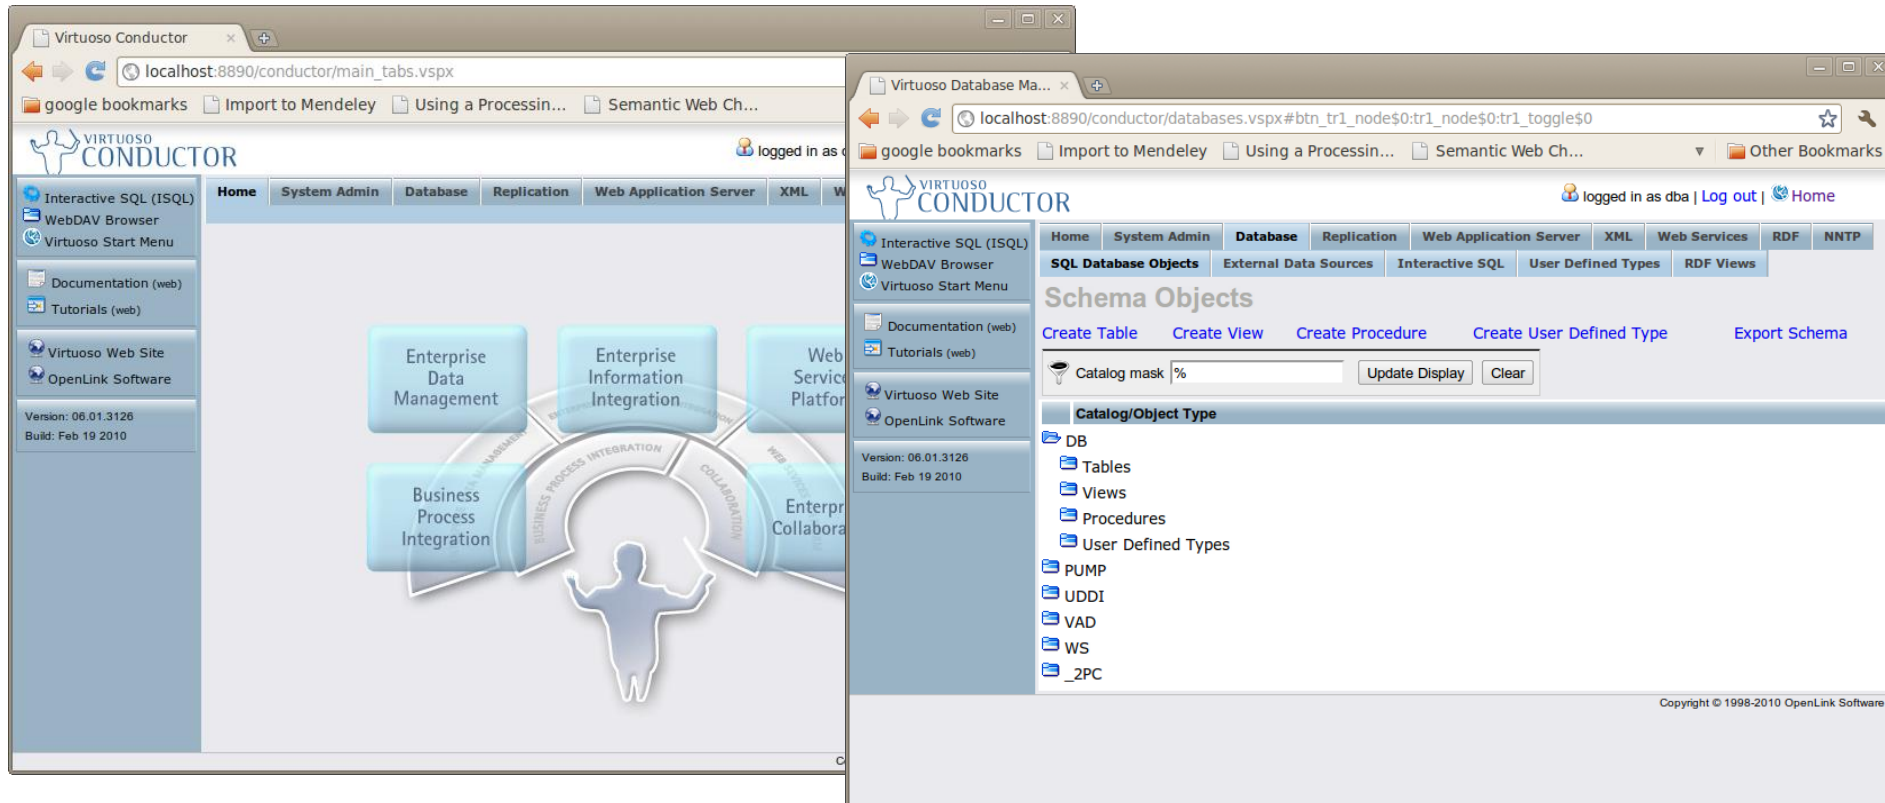


Exposing sensor data as triples (3)

- Using Virtuoso, RDF Data can also be accessible via
 - ODBC/JDBC
 - ADO.NET (Entity Frameworks compatible)
 - OLE DB
 - XMLA (XML for Analysis) data providers / drivers
- Using the “Sponger” RDF-izer, RDF data can be extracted from non-RDF sources (e.g. with XSLT)

Publishing RDF using Virtuoso (1)

- Conductor: a GUI for server administration
- Virtuoso can be used as a DBMS



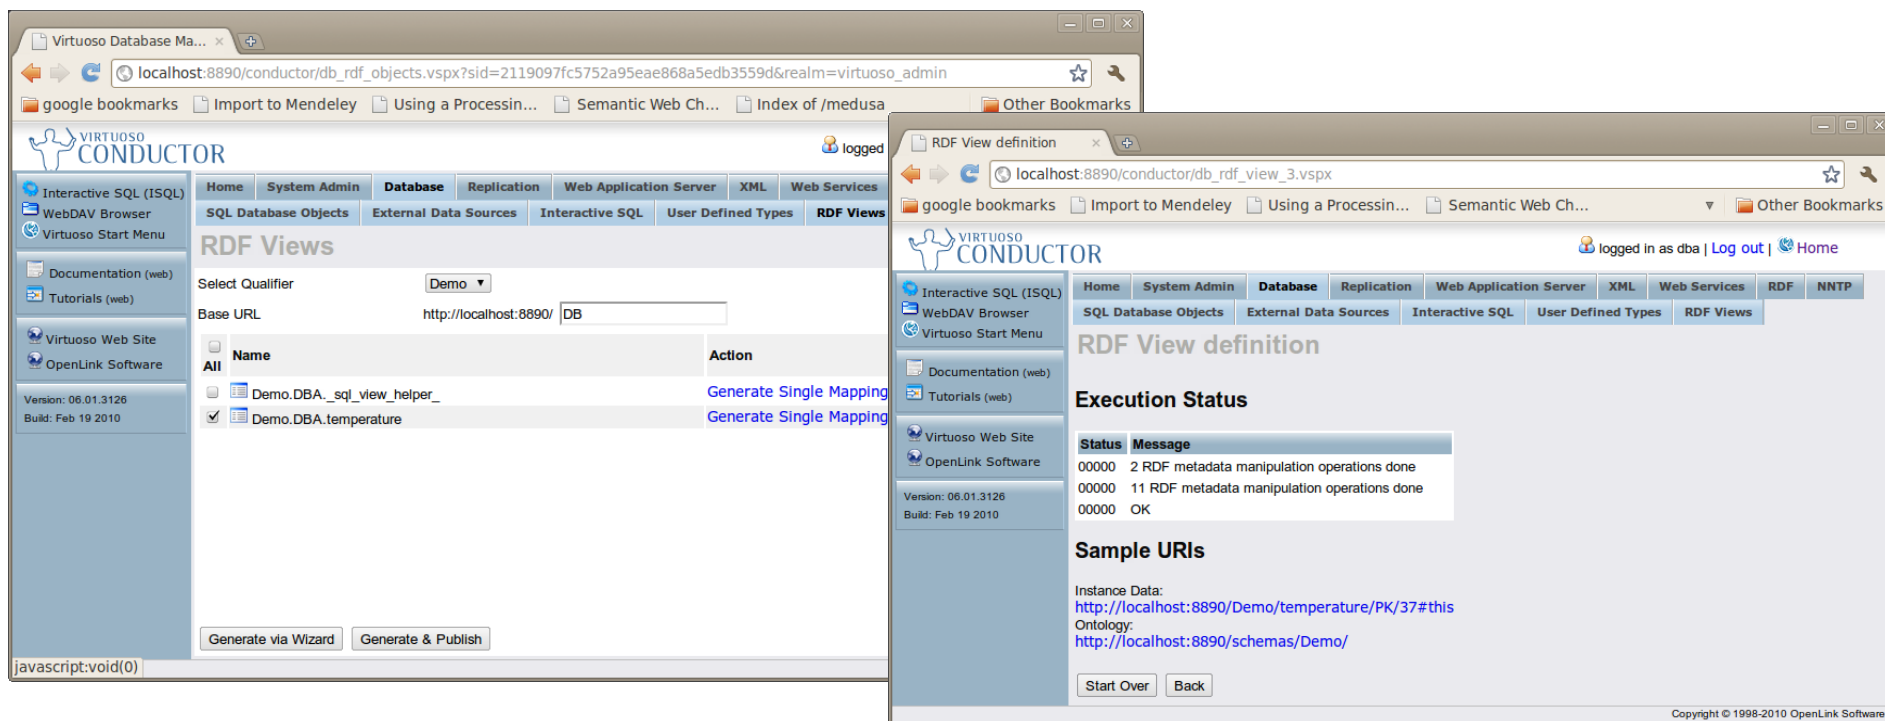
The image displays two screenshots of the Virtuoso Conductor web interface, illustrating its role as a GUI for server administration and its use as a DBMS.

Left Screenshot (main_tabs.vsp): This view shows the main navigation menu with tabs for Home, System Admin, Database, Replication, Web Application Server, XML, and Web Services. The central area features a diagram illustrating the integration of various components: Enterprise Data Management, Enterprise Information Integration, Web Service Platform, Business Process Integration, and Enterprise Collaboration. The interface is logged in as 'dba'.

Right Screenshot (databases.vsp): This view shows the 'Schema Objects' section, which is used for managing the database. It includes a 'Catalog mask' field and a list of objects under the 'Catalog/Object Type' section. The objects listed are: DB, Tables, Views, Procedures, User Defined Types, PUMP, UDDI, VAD, WS, and _2PC. The interface is logged in as 'dba'.

Publishing RDF using Virtuoso (2)

- Can be combined with GSN to process sensor data streams and export them as RDF
- Create RDF Views over the relational data



The image displays two screenshots of the Virtuoso Conductor web interface, illustrating the process of publishing RDF using Virtuoso.

Left Screenshot: RDF Views Management

The interface shows the 'Virtuoso Database Management System' (Virtuoso Conductor) with the 'RDF Views' tab selected. The 'RDF Views' section displays a table of existing views:

Name	Action
Demo.DBA__sql_view_helper_	Generate Single Mapping
Demo.DBA.temperature	Generate Single Mapping

Buttons at the bottom include 'Generate via Wizard' and 'Generate & Publish'.

Right Screenshot: RDF View definition

The interface shows the 'RDF View definition' page. The 'Execution Status' section displays the following information:

Status	Message
00000	2 RDF metadata manipulation operations done
00000	11 RDF metadata manipulation operations done
00000	OK

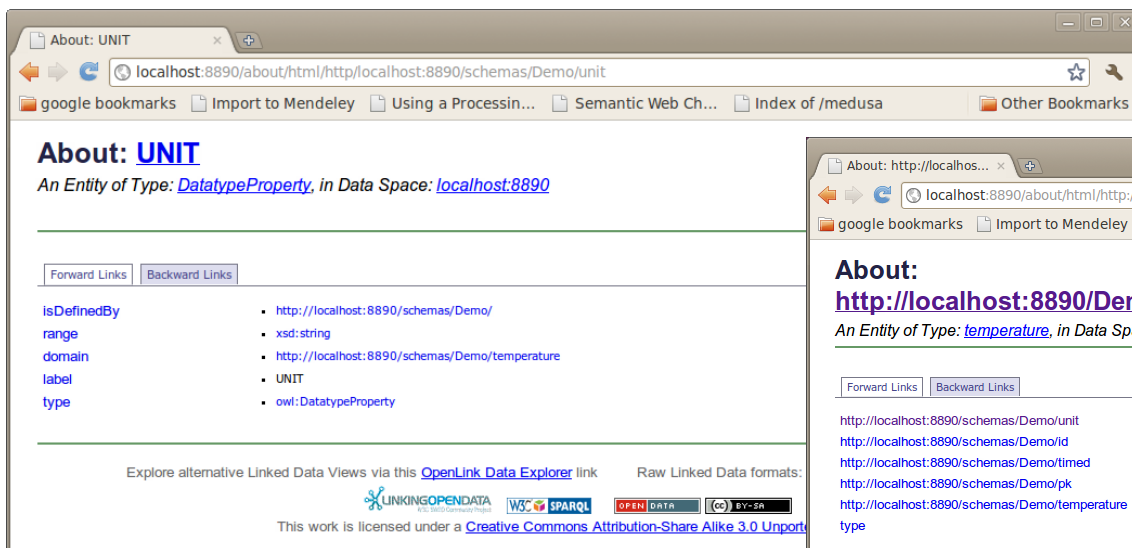
The 'Sample URIs' section displays the following information:

Instance Data: <http://localhost:8890/Demo/temperature/PK/37#this>
Ontology: <http://localhost:8890/schemas/Demo/>

Buttons at the bottom include 'Start Over' and 'Back'.

Publishing RDF using Virtuoso (3)

- Browseable repository
- A URI for every resource
- Example: Measurement URI



About: UNIT

localhost:8890/about/html/http://localhost:8890/schemas/Demo/unit

google bookmarks Import to Mendeley Using a Processin... Semantic Web Ch... Index of /medusa Other Bookmarks

About: **UNIT**

An Entity of Type: [DatatypeProperty](#), in Data Space: [localhost:8890](#)

Forward Links Backward Links

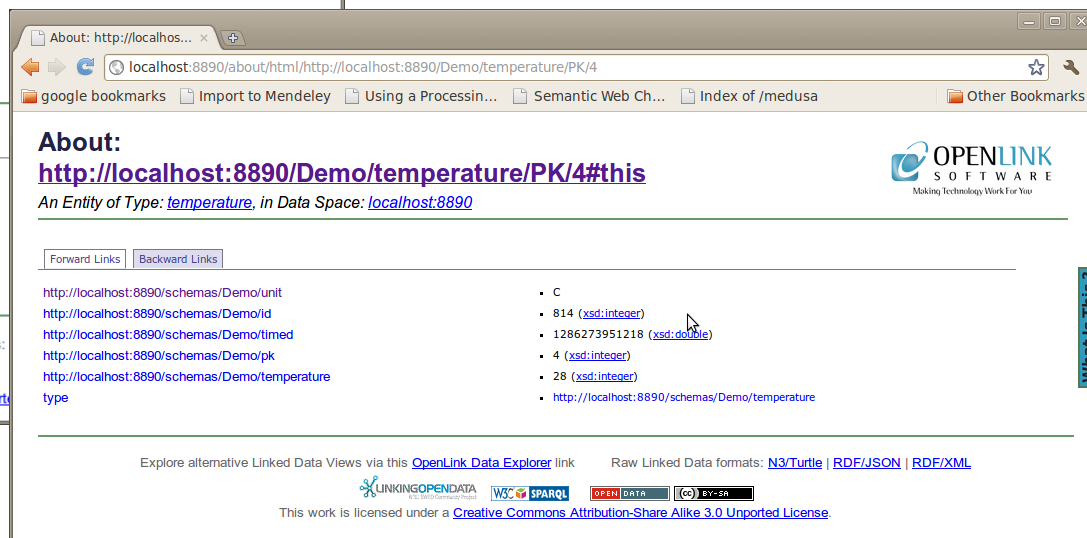
isDefinedBy

- [http://localhost:8890/schemas/Demo/](#)
- [xsd:string](#)
- [http://localhost:8890/schemas/Demo/temperature](#)
- [UNIT](#)
- [owl:DatatypeProperty](#)

Explore alternative Linked Data Views via this [OpenLink Data Explorer](#) link Raw Linked Data formats:

LINKINGOPENDATA W3C SPARQL OPEN DATA CC BY-SA

This work is licensed under a [Creative Commons Attribution-Share Alike 3.0 Unported](#)



About: http://localhost:8890/Demo/temperature/PK/4#this

localhost:8890/about/html/http://localhost:8890/Demo/temperature/PK/4

google bookmarks Import to Mendeley Using a Processin... Semantic Web Ch... Index of /medusa Other Bookmarks

About: **http://localhost:8890/Demo/temperature/PK/4#this**

An Entity of Type: [temperature](#), in Data Space: [localhost:8890](#)

Forward Links Backward Links

http://localhost:8890/schemas/Demo/unit

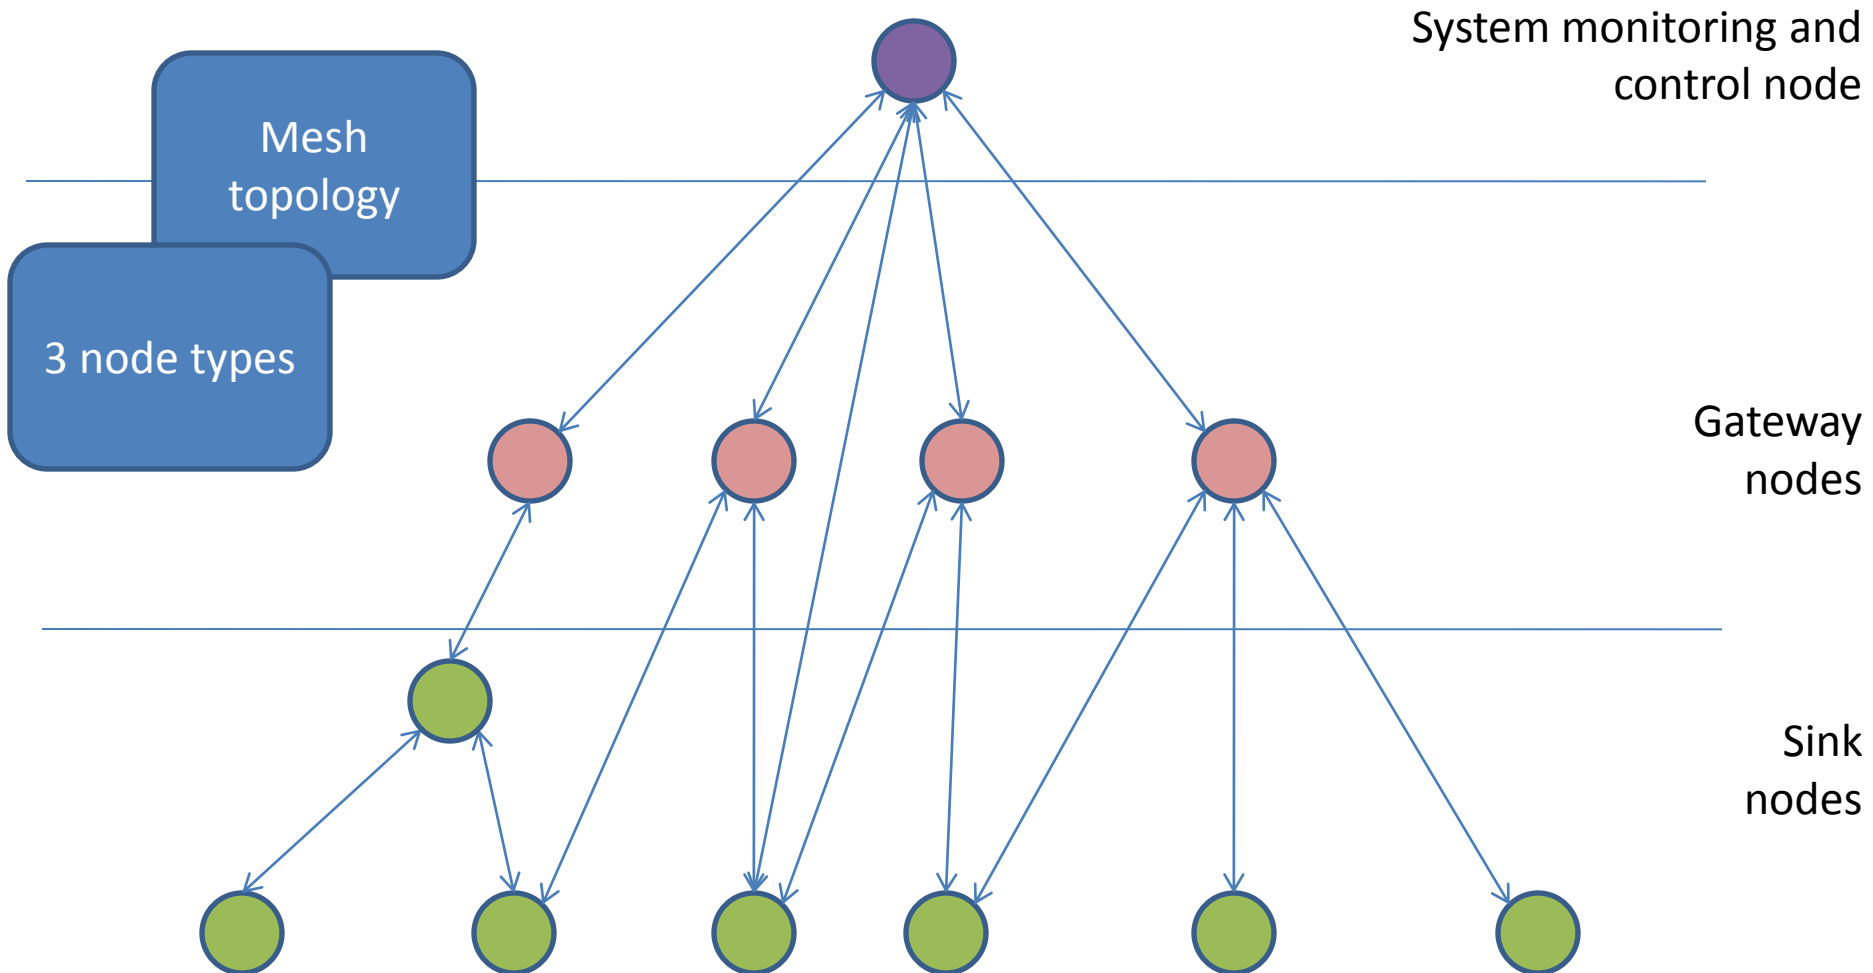
- C
- 814 ([xsd:integer](#))
- 1286273951218 ([xsd:double](#))
- 4 ([xsd:integer](#))
- 28 ([xsd:integer](#))
- [http://localhost:8890/schemas/Demo/temperature type](#)

Explore alternative Linked Data Views via this [OpenLink Data Explorer](#) link Raw Linked Data formats: [N3/Turtle](#) | [RDF/JSON](#) | [RDF/XML](#)

LINKINGOPENDATA W3C SPARQL OPEN DATA CC BY-SA

This work is licensed under a [Creative Commons Attribution-Share Alike 3.0 Unported License](#)

Semantic Sensor Network Example (1)



Semantic Sensor Network Example (2)

- Sink node
 - Operation relies on a relational database
 - Limited historical data
 - Keep a “sliding window”
 - Based on time or tuples
 - Do not have semantic capabilities
 - One Database per Sink node



Semantic Sensor Network Example (3)

- Gateway Node
 - Operation relies on a semantically-enabled knowledge base
 - Supported by inference procedures
 - Maintains historical/archived information
 - One Knowledge Base per Gateway node
 - Appropriate for Higher Level Fusion (e.g. threats, events)

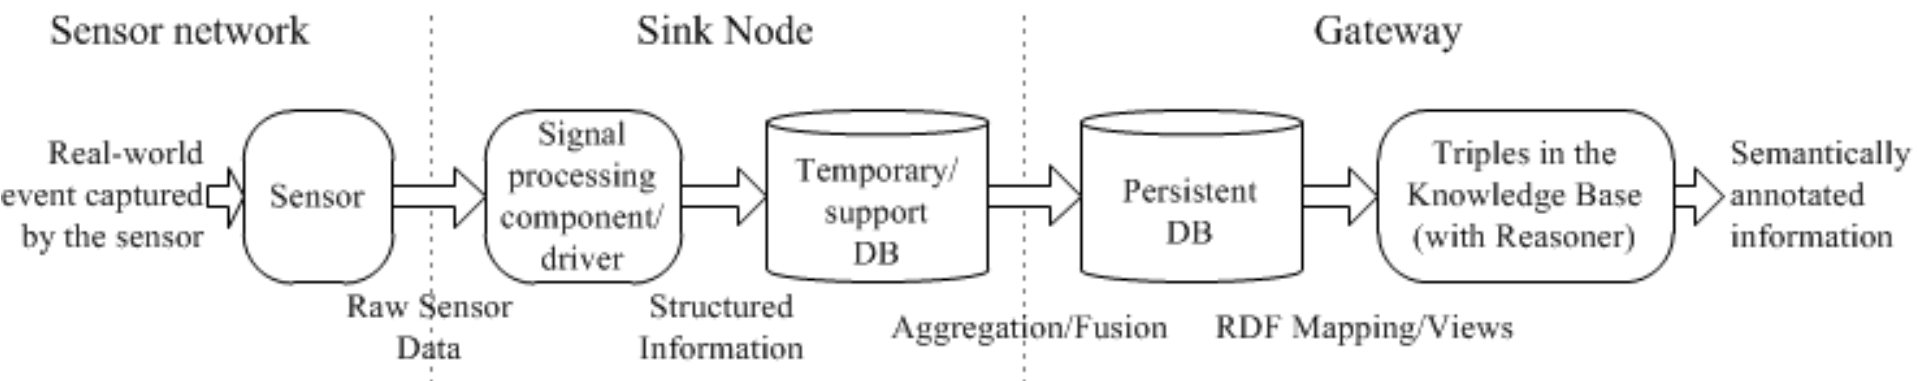


Semantic Sensor Network Example (4)

- Central Control node
 - Monitors and controls the network
 - Provides system-wide services
 - E.g. directory services, secure authentication
 - Can store its view over the network for intelligence extraction

Information Flow

- An Information flow example in a decentralized Semantic Sensor Network





Lecture Outline

- Introduction
- Sensor data
- Semantic web
- Context-awareness
- The GSN middleware
- Exposing sensor data as triples
- Semantic Integration using SPARQL

Semantic integration using SPARQL (1)

- SPARQL: An SQL-like language for querying RDF graphs
- SELECT-FROM-WHERE syntax
- WHERE conditions are triple patterns
- `SELECT ?x ?y ?z`

`WHERE`

`{ ?x ?y ?z }`

returns all the triples in the graph



Semantic integration using SPARQL (2)

- XML over HTTP (RESTful approach)
 - http://demo.openlinksw.com/sparql?default-graph-uri=urn:lsid:ubio.org:namebank:11815&should-sponge=soft&query=SELECT+*+WHERE+{?s+?p+?o}&format=text/html
- No create/update/delete capabilities

Semantic integration using SPARQL (3)

- SPARQL queries can be named and stored
 - A query named *sparql-demo* listens to:
<http://localhost:8890/DAV/sparql-demo>
- Can return results over HTTP (XML by default)
- MIME type of the RDF data
 - 'rdf+xml' (default) | 'n3' | 'turtle' | 'ttl'



Semantic integration using SPARQL (4)

- SPARQL results example in RDF/XML

```
<ROOT>
<rdf:RDF xmlns:rdf="http://www.w3.org/1999/02/22-rdf-syntax-ns#" xmlns:rs="http://www.w3.org/2005/sparql-results#"
xmlns:xsd="http://www.w3.org/2001/XMLSchema#">
  <rs:results rdf:nodeID="rset">
    <rs:result rdf:nodeID="sol193">
      <rs:binding rdf:nodeID="sol193-0" rs:name="x"><rs:value rdf:resource="http://localhost:8890/Demo/temperature/PK/4#this"/></rs:binding>
      <rs:binding rdf:nodeID="sol193-1" rs:name="y"><rs:value rdf:resource="http://www.w3.org/1999/02/22-rdf-syntax-ns#type"/></rs:binding>
      <rs:binding rdf:nodeID="sol193-2" rs:name="z"><rs:value rdf:resource="http://localhost:8890/schemas/Demo/temperature"/></rs:binding>
    </rs:result>
  ...
</rs:results>
</rdf:RDF>
</ROOT>
```



Thank you!

Questions?